

Designing Kubernetes Implementation

Overall Design (Level 0)

- Not a single day activity
- Multiple workphases (3 weeks document)
- Start from zero

- List requirements (e.g., UI, Payments, Transactions, Reports)
- Identify teams working on microservices (helps define dependencies)
- Categorize microservices by business criticality
 - Low critical
 - Medium critical
 - Important
 - Critical

- Low critical
- Medium critical
- Important
- Critical
- CPU utilization
- Memory utilization
- Disk utilization
- Services: Prometheus, Grafana, AWS Billing/CloudWatch, Azure Dashboard
- Helps define Kubernetes cluster size (CPU/Memory/Storage)

Level Zero: Requirement Gathering

- Attack primary task
- Don't create cluster directly
- Understand requirements
- Key Activities

- List microservices (e.g., UI, Payments, Transactions, Reports)
- Identify teams working on microservices (helps define dependencies)
- Categorize microservices by business criticality
 - Low critical
 - Medium critical
 - Important
 - Critical
- Identify existing resource utilization (if available)
 - CPU utilization
 - Memory utilization
 - Disk utilization
 - Services: Prometheus, Grafana, AWS Billing/CloudWatch, Azure Dashboard
 - Helps define Kubernetes cluster size (CPU/Memory/Storage)

- Calculate Cost
 - Current cost (VMs)
 - Estimated Kubernetes cost (CPU, RAM, SSD, S3)
 - Reduce cost vs performance
- Document findings (e.g., "Spreadsheets")

- Communications
 - Show document with management
 - Show document with development teams
 - Discuss the plan

Level One: Proof of Concept (POC)

- Don't jump to Deploying Prod clusters
- Show implementation with an Kubernetes
- Identify potential issues
- Activities
- Typical duration: ~30 days (or 2-3 weeks)

- Set up 15-20 representative apps (sets of microservices)
- Create small Kubernetes cluster (e.g., 1 control plane, 3 data plane nodes, 8 CPU/16 GB RAM on AWS)
- Create deployment manifests (Deployment, Service, Ingress, ConfigMap)
- Push configuration of stateful stateless apps (DB, app, caching, queueing)
- Provision Kubernetes Service
- Provision Ingress (e.g., Cloud provider's controller like ALB)
- Test traffic
- QA/UAT on main data
- Test timeout issues (e.g., crash loop back off)
- Configure metrics/monitoring probes

- Start actual implementation
- Use data from Level 0 & 1
- Keep lower Kubernetes environment
- Cluster Sizing & Nodes

- Use resource utilization data (Level 0)
- Control plane: At least 3 nodes (if running on-premise)
- Data plane: At least 3 nodes
- Advance CPU/Memory - required
- Keep number of nodes manageable for upgrade

Level Two: Dev Kubernetes Cluster Setup

Prerequisites

- Logical isolation
- Basic isolation: Per team
- Enables RBAC (Role-Based Access Control)
- Integrate RBAC with IAM (e.g., OIDC)

Resource Management within Cluster

- Resource Quotas (per Namespace)
 - Maximum resources a namespace can utilize
 - Prevent one namespace from using all resources
- Limit Ranges (per Pod)
 - Enforced pods from using resources
 - CPU limits
 - Memory limits
- Requests and Limits (per Pod): Used to restrict pod resource usage

Level Three: Staging/High Environment Setup

- Build upon Dev setup
- Required for stable testing environment
- Test production related scenarios
- Dev environment is often available (developer's request)
- Implementation Approaches (Two Popular Steps)
- Number of environments depends on business utility (Dev, QA, Staging, UAT, Pre-prod, Prod)
- Follow same setup approach: Resource, Prerequisite, Resource Quotas, Limits
- Recommendation: Staging environment with extra resources behind to protect

- Start Dev Cluster with QA
 - Use the same Dev cluster (shared IP)
 - Network restriction
 - Aggregate using Namespace (e.g., dev-payments, stage-payments)
 - Requires strict RBAC management
 - Complexity in RBAC and IAM
- Create Separate Staging Cluster
 - Replicate Dev environment setup
 - Separate dev and staging clusters
 - Clear isolation
 - Prefered approach if budget allows

Level Four: Production Environment Setup

- Setup to mirror after Dev/Staging
- Key Considerations
 - Multi Availability Zone (Multi AZ)
 - Relatively required for high availability
 - Hardware nodes across multiple AZs (data centers)
 - Distribute pod/replicas across multiple AZs
 - Use Kubernetes Topology Spread Constraints
 - Multi Region Setup (if required for global customers)
 - Address network latency
 - Higher availability scaling (Level 5)
- Resource Management: More limited, or use Autoscaling (if supported) or manual scaling with more resources

- Prerequisites before Scaling (Level 5)
 - Good Observability Stack (Prometheus, Grafana)
 - Control Loops/Random Probes
 - Control pod resource assignment
- Typically involves high availability and multi-region setup
 - Deploy pods across different clusters
- High Availability Setup (Multi-Cluster)
 - Separate for each cluster
 - Inter-face with Global Load Balancer
 - AWS Load Balancer
 - DNS based load balancing (e.g., Route 53)
 - Route requests based on origin region

Further Levels (Beyond Level 5)

- Container
- Linux distro
- Security
- Networking
- OS/Kernel
- Cloud